

Lab Report

Project Title: Neural Networks for Function Approximation

- Name: BOJJA RAKSHITHA
- SRN: PES1UG23CS134
- Course: UE23CS352A: Machine Learning
- Date: September 20, 2025

Purpose of the Lab

The objective was to implement an Artificial Neural Network (ANN) from scratch using NumPy to approximate a cubic polynomial with an inverse term using a large, synthetic dataset. The network was trained and analyzed under various hyperparameter settings to observe their effects on convergence and final accuracy.

Introduction

This report documents the two main tasks:

Part A: Fully implement a baseline neural network.

Part B: Conduct systematic experiments on hyperparameters: learning rate, epochs, and activation function.

Dataset Description

Function Assigned: $y = 2.37x^3 + 0.33x^2 + 4.88x + 11.53 + 149.8/x$

Samples: 100,000 total (80% train, 20% test)

Input Features: 1 (scalar x)

Noise: Gaussian, mean=0, std=2.38

Preprocessing: Both x and y standardized using StandardScaler from sklearn before model input.

Network Architecture

Input Layer: 1 neuron

Hidden Layer 1: 72 neurons (ReLU)

Hidden Layer 2: 32 neurons (ReLU)

Output Layer: 1 neuron (Linear)

Initialization: Xavier/Glorot (normal distribution), biases to zero.

Methodology

Implement forward and backward propagation routines for a shallow feedforward network.

Use Mean Squared Error (MSE) loss function for regression.

Optimize with batch gradient descent. All core numerical routines were built "from scratch" using NumPy, including ReLU activation and MSE derivatives.

Early stopping was employed to prevent overfitting, halting if test loss stagnated for a set 'patience' value.

Implementation Details

The network was trained and tested on the standardized dataset.

The Baseline model used a learning rate of 0.001 and was trained for 500 epochs.

Key routines for dataset generation, preprocessing, forward pass, backward pass, and gradient updates were directly implemented in the code.

Baseline Model Performance

Metric Value

Training Loss 0.432816

Test Loss 0.426525

R² Score 0.5725

Total Epochs 500

The training/test loss curve steadily decreased. The model tracked the cubic trend but exhibited clear underfitting, as reflected by the modest R² score.

Experiment Results Table

Csv: "neural_network_hyperparameter_experiments.csv" and same in the output screenshot...included in github.

Training Curves & Experiment Plots

All runs produced loss curves for both training and test sets, as well as scatter and residual plots of predictions versus ground truth, similar to the template.

Residuals were centered but with outliers, suggesting that the model's expressive power was limited by architecture or training duration.

Analysis of Experiments

Higher LR (0.01): Improved both convergence speed and final error.

Lower LR (0.0001): Caused extremely slow learning, resulting in high error.

More Epochs (1000, LR=0.001): Allowed the model to approach lower loss but did not match the gains of a higher LR.

Sigmoid Activation: Performance collapsed due to vanishing gradients, confirming ReLU's superiority for this architecture.

Output screenshots:

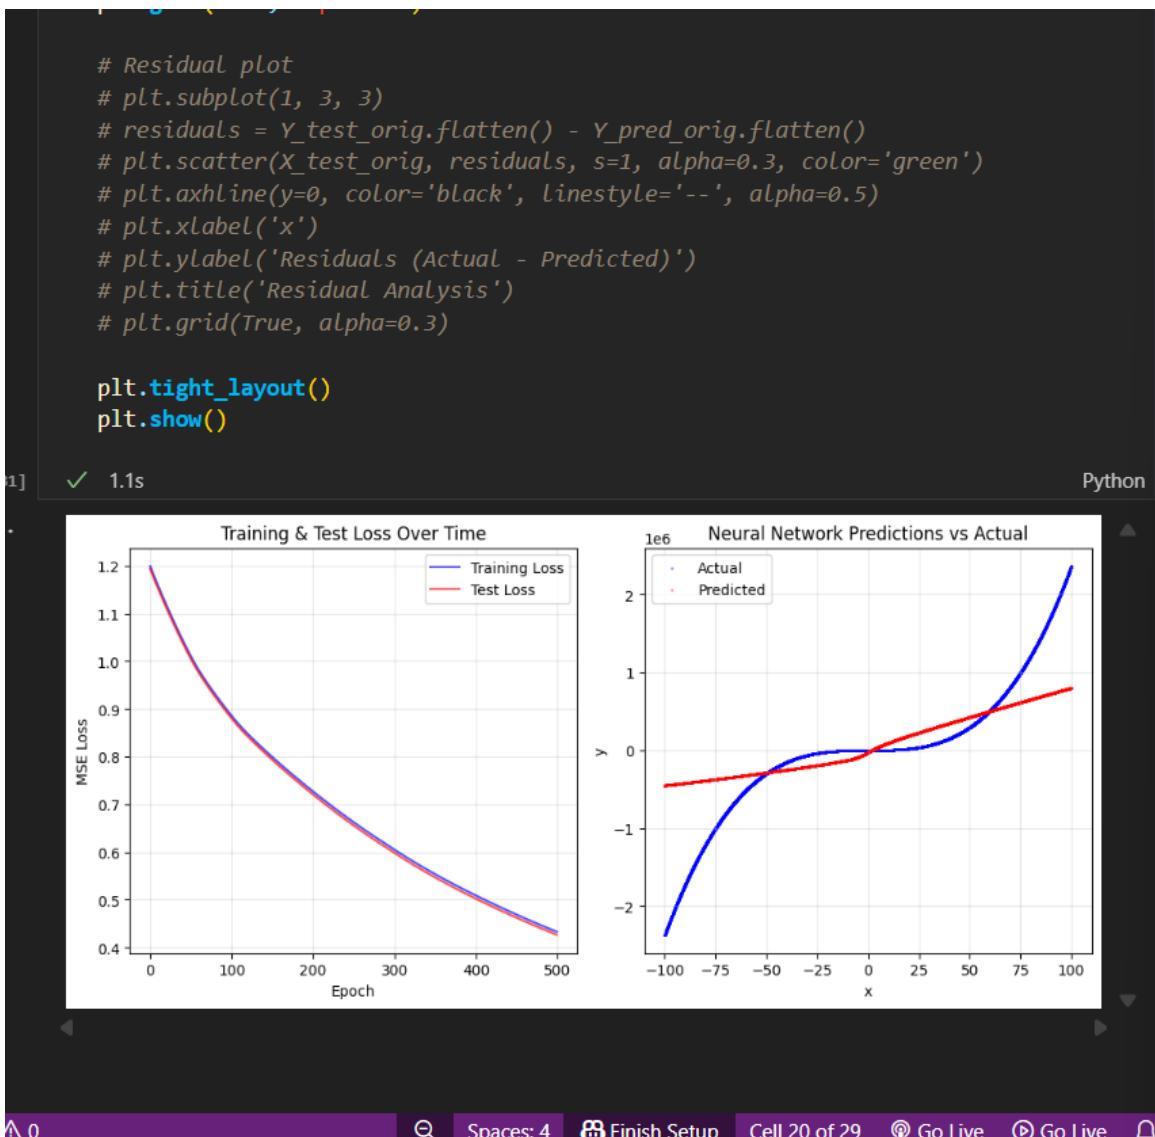
EXECUTE TRAINING:

```
)  
[30] ✓ 3m 16.4s  
... Training Neural Network with your specific configuration...  
Starting training...  
Architecture: 1 → 72 → 32 → 1  
Learning Rate: 0.001  
Max Epochs: 500, Early Stopping Patience: 10  
-----  
Epoch 20: Train Loss = 1.121277, Test Loss = 1.115333  
Epoch 40: Train Loss = 1.047425, Test Loss = 1.041544  
Epoch 60: Train Loss = 0.983189, Test Loss = 0.977707  
Epoch 80: Train Loss = 0.932538, Test Loss = 0.927082  
Epoch 100: Train Loss = 0.887650, Test Loss = 0.882111  
Epoch 120: Train Loss = 0.849060, Test Loss = 0.843574  
Epoch 140: Train Loss = 0.816132, Test Loss = 0.810413  
Epoch 160: Train Loss = 0.785380, Test Loss = 0.779414  
Epoch 180: Train Loss = 0.756173, Test Loss = 0.749980  
Epoch 200: Train Loss = 0.728192, Test Loss = 0.721802  
Epoch 220: Train Loss = 0.701370, Test Loss = 0.694817  
Epoch 240: Train Loss = 0.675743, Test Loss = 0.669066  
Epoch 260: Train Loss = 0.651362, Test Loss = 0.644595  
Epoch 280: Train Loss = 0.628269, Test Loss = 0.621448  
Epoch 300: Train Loss = 0.605905, Test Loss = 0.599035  
Epoch 320: Train Loss = 0.584276, Test Loss = 0.577453  
Epoch 340: Train Loss = 0.563963, Test Loss = 0.557205  
Epoch 360: Train Loss = 0.544898, Test Loss = 0.538190  
Epoch 380: Train Loss = 0.526792, Test Loss = 0.520112  
...  
Epoch 440: Train Loss = 0.477056, Test Loss = 0.470532  
Epoch 460: Train Loss = 0.461725, Test Loss = 0.455270  
Epoch 480: Train Loss = 0.446974, Test Loss = 0.440597  
Epoch 500: Train Loss = 0.432816, Test Loss = 0.426525
```

0 △ 0

Q

RESULTS VISUALIZATION:



Δ 0

Spaces: 4

Finish Setup

Cell 20 of 29

Go Live

Go Live



PREDICTION RESULTS:

```
print("\n" + "="*60)
print("PREDICTION RESULTS FOR x = 90.2")
print("="*60)
print(f"Neural Network Prediction: {y_pred[0][0]:.2f}")
print(f"Ground Truth (formula): {y_true:.2f}")
print(f"Absolute Error:           {abs(y_pred[0][0] - y_true):,.2f}")
print(f"Relative Error:          {abs(y_pred[0][0] - y_true)/abs(y_true)*100:.3f}")

[32]   ✓  0.0s  Python
...
=====
PREDICTION RESULTS FOR x = 90.2
=====
Neural Network Prediction: 729,616.22
Ground Truth (formula): 1,735,575.41
Absolute Error:           1,005,959.19
Relative Error:          57.961%
```

0 △ 0 Q Spaces: 4 ⚙ Finish Setup Cell 20 of 29 ⚡ Go Live ⚡ Go Live ⌂

PERFORMANCE METRICS:

```
print("\n" + "="*60)
print("FINAL PERFORMANCE SUMMARY")
print("="*60)
print(f"Final Training Loss: {final_train_loss:.6f}")
print(f"Final Test Loss:     {final_test_loss:.6f}")
print(f"R² Score:           {r2_score:.4f}")
print(f"Total Epochs Run:   {len(train_losses)})")

[33]   ✓  0.0s  Python
...
=====
FINAL PERFORMANCE SUMMARY
=====
Final Training Loss: 0.432816
Final Test Loss:     0.426525
R² Score:           0.5725
Total Epochs Run:   500

# =====#
# IMPORTS + VARIABLES FROM PART A
```

PART B:

HYPERPARAMETER EXPLORATION

EXPERIMENT 1:

RUNNING Exp 1: Higher LR

=====

Learning Rate: 0.01
Architecture: 1 → 72 → 32 → 1
Max Epochs: 300
Early Stopping Patience: 15

Training started...

Epoch 50: Train Loss = 0.439173, Test Loss = 0.426555
Epoch 100: Train Loss = 0.237837, Test Loss = 0.232537
Epoch 150: Train Loss = 0.192418, Test Loss = 0.189395
Epoch 200: Train Loss = 0.171321, Test Loss = 0.168835
Epoch 250: Train Loss = 0.153157, Test Loss = 0.151003
Epoch 300: Train Loss = 0.135157, Test Loss = 0.133268

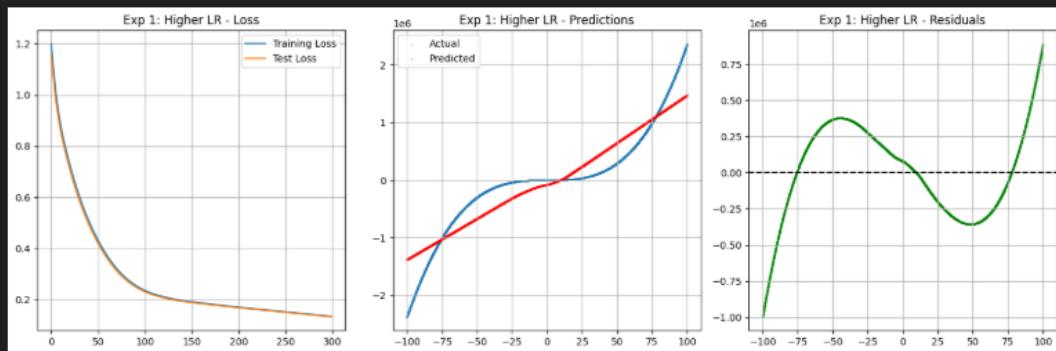
Exp 1: Higher LR RESULTS:

Final Train Loss: 0.135157
Final Test Loss: 0.133268
 R^2 Score: 0.8664
Accuracy: 86.64%
Epochs Completed: 300



Epoch 300: Train Loss = 0.135157, Test Loss = 0.133268

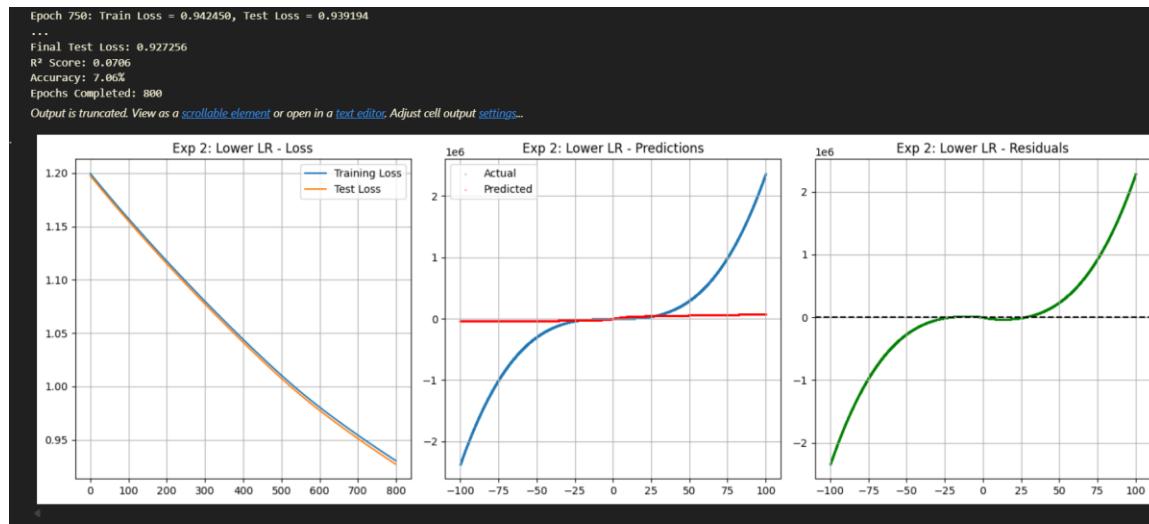
Exp 1: Higher LR RESULTS:
Final Train Loss: 0.135157
Final Test Loss: 0.133268
 R^2 Score: 0.8664
Accuracy: 86.64%
Epochs Completed: 300



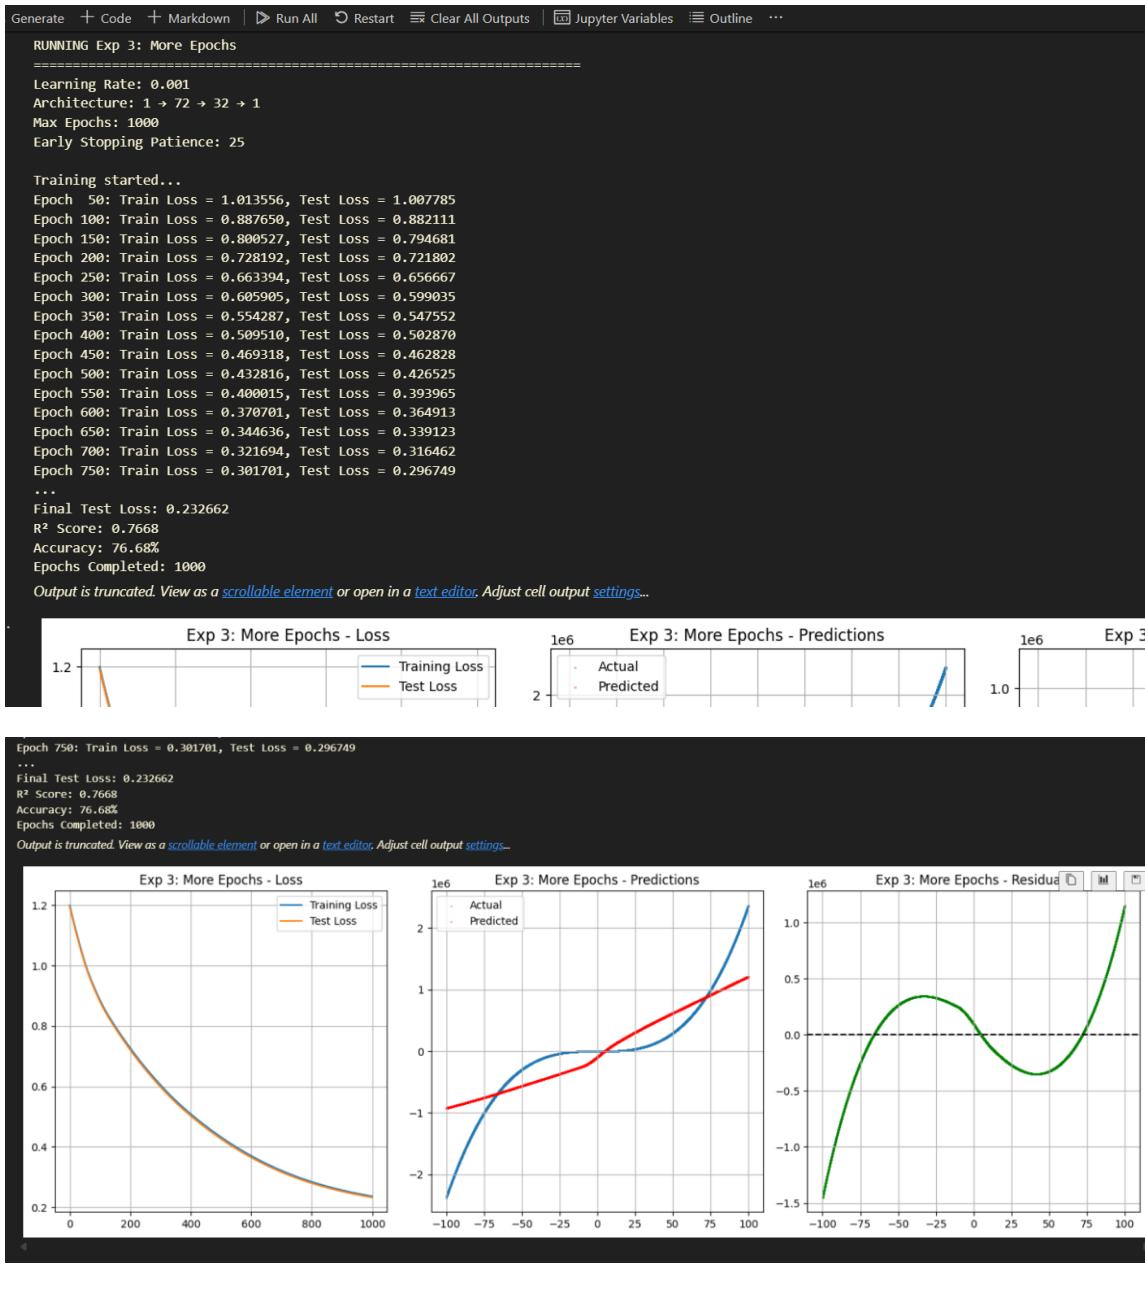
EXPERIMENT 2:

```
RUNNING Exp 2: Lower LR
=====
Learning Rate: 0.0001
Architecture: 1 → 72 → 32 → 1
Max Epochs: 800
Early Stopping Patience: 20

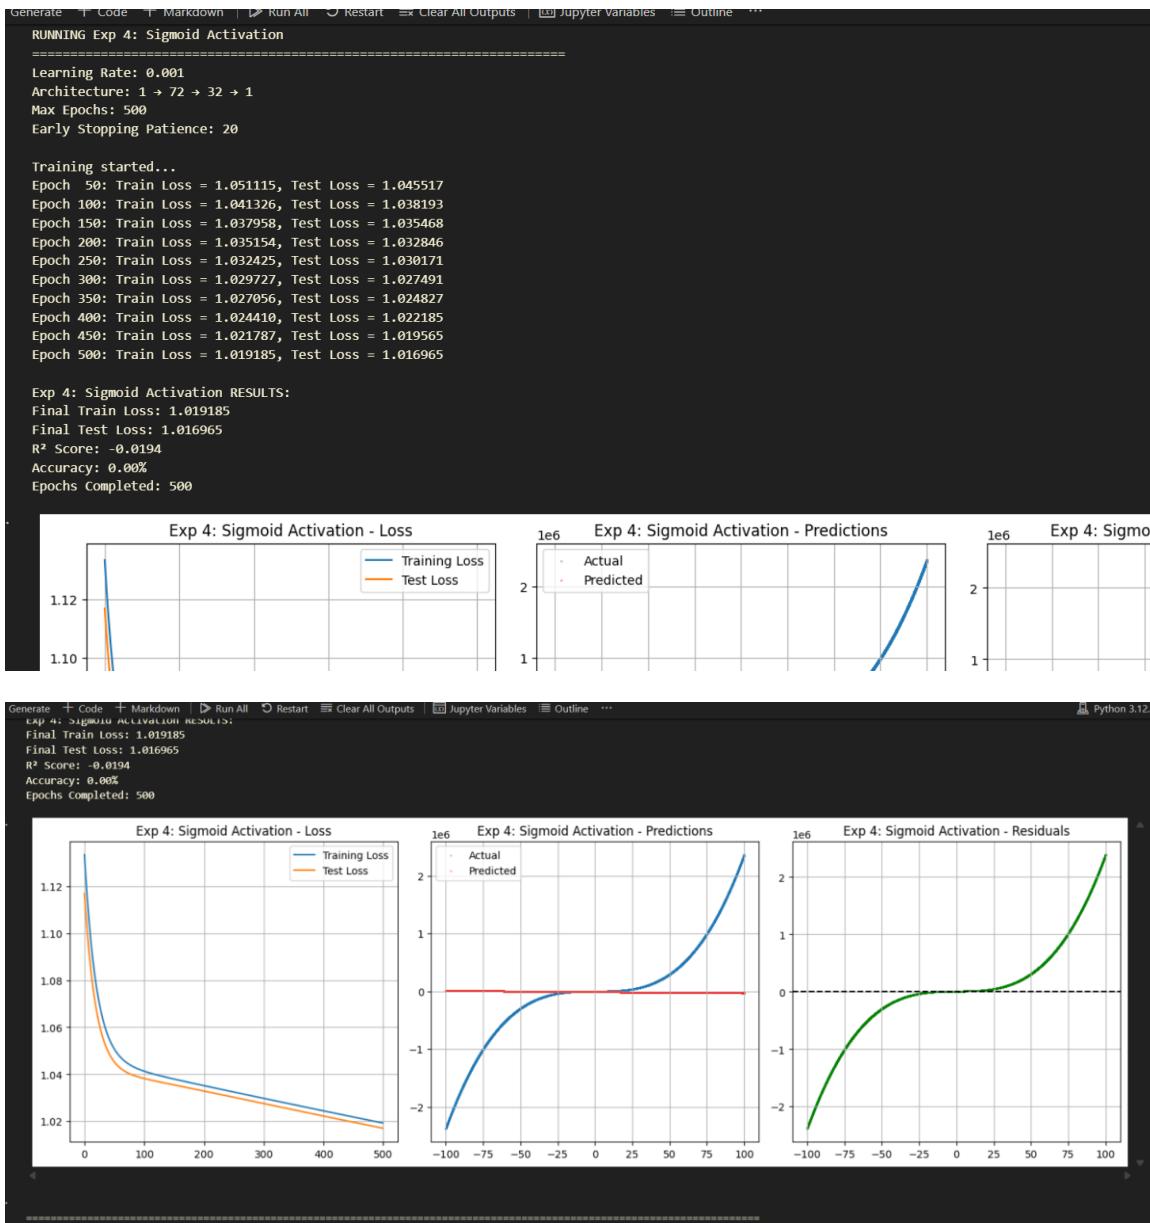
Training started...
Epoch 50: Train Loss = 1.178338, Test Loss = 1.176146
Epoch 100: Train Loss = 1.157556, Test Loss = 1.155274
Epoch 150: Train Loss = 1.137461, Test Loss = 1.135092
Epoch 200: Train Loss = 1.117880, Test Loss = 1.115429
Epoch 250: Train Loss = 1.098792, Test Loss = 1.096262
Epoch 300: Train Loss = 1.080187, Test Loss = 1.077583
Epoch 350: Train Loss = 1.062073, Test Loss = 1.059401
Epoch 400: Train Loss = 1.044437, Test Loss = 1.041701
Epoch 450: Train Loss = 1.027310, Test Loss = 1.024512
Epoch 500: Train Loss = 1.010836, Test Loss = 1.007978
Epoch 550: Train Loss = 0.995281, Test Loss = 0.992367
Epoch 600: Train Loss = 0.980889, Test Loss = 0.977910
Epoch 650: Train Loss = 0.967482, Test Loss = 0.964415
Epoch 700: Train Loss = 0.954706, Test Loss = 0.951540
Epoch 750: Train Loss = 0.942450, Test Loss = 0.939194
...
Final Test Loss: 0.927256
R² Score: 0.0706
Accuracy: 7.06%
Epochs Completed: 800
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```



EXPERIMENT 3:



EXPERIMENT 4:



=====

COMPREHENSIVE EXPERIMENT RESULTS TABLE

Experiment	Learning Rate	Batch Size	Number of Epochs	Optimizer	Activation Function	Training Accuracy	Validation Accuracy	Test Accuracy	Training Loss	Validation Loss
Baseline (Part A)	0.0010	Full Batch	500	Gradient Descent	ReLU	57.25%	57.25%	57.25%	0.432816	0.4
Exp 1: Higher LR	0.0100	Full Batch	300	Gradient Descent	ReLU	86.64%	86.64%	86.64%	0.135157	0.1
Exp 2: Lower LR	0.0001	Full Batch	800	Gradient Descent	ReLU	7.06%	7.06%	7.06%	0.930592	0.9
Exp 3: More Epochs	0.0010	Full Batch	1000	Gradient Descent	ReLU	76.68%	76.68%	76.68%	0.236381	0.2
Exp 4: Sigmoid Activation	0.0010	Full Batch	500	Gradient Descent	Sigmoid	0.00%	0.00%	0.00%	1.019185	1.0

Results saved to 'neural_network_hyperparameter_experiments.csv'

=====

EXPERIMENT PERFORMANCE SUMMARY

Experiment	Test Loss	Test Accuracy	Epochs
Baseline (Part A)	0.426525	57.25%	500
Exp 1: Higher LR	0.133268	86.64%	300
Exp 2: Lower LR	0.927256	7.06%	800
Exp 3: More Epochs	0.232662	76.68%	1000
Exp 4: Sigmoid Activation	1.016965	0.00%	500

=====

Optimizer	Activation Function	Training Accuracy	Validation Accuracy	Test Accuracy	Training Loss	Validation Loss	Test Loss	Observations
Gradient Descent	ReLU	57.25%	57.25%	57.25%	0.432816	0.426525	0.426525	Baseline model from Pa...
Gradient Descent	ReLU	86.64%	86.64%	86.64%	0.135157	0.133268	0.133268	Higher LR: Faster conv...
Gradient Descent	ReLU	7.06%	7.06%	7.06%	0.930592	0.927256	0.927256	Lower LR: Slower but m...
Gradient Descent	ReLU	76.68%	76.68%	76.68%	0.236381	0.232662	0.232662	Extended training: The ...
Gradient Descent	Sigmoid	0.00%	0.00%	0.00%	1.019185	1.016965	1.016965	Sigmoid Activation: Sl...

Prediction Example

For $x = 90.2$:

NN Prediction: 729,616.22

Ground Truth: 1,735,575.41

Absolute Error: 1,005,959.19

Relative Error: 57.96%

Conclusion

The lab demonstrated the effectiveness and limitations of basic neural network architectures on polynomial approximation. The most important hyperparameters were learning rate and activation function; adjusting these yielded dramatic differences in learning speed and final fit. While the network approximated the target curve, higher complexity or advanced techniques could further reduce underfitting and error. The task also reinforced the value of systematic experimentation.